

Project Repository

Simulator for Reinforcement Learning-based Fleet Management

GitHub Repository

<https://github.com/illidanlab/Simulator>

A large-scale ride-hailing simulator supporting IDQN, contextual DQN, contextual A2C, and LP-cA2C algorithms.

algorithm/alg_utility.py |

Purpose: Provides **helper functions** and **math utilities** that all RL algorithms (IDQN, cDQN, cA2C, LP-cA2C) use inside the simulator.

Key Components

- ▶ **Probability Distribution Classes (Pd, DiagGaussianPd)**
Used to sample actions from a **Gaussian policy** (for Actor-Critic). Computes log-probabilities, entropy, sampling.
- ▶ **Reward Normalization**
Makes rewards stable by converting them to **zero-mean & unit-variance**.
- ▶ **Continuous Quadratic Knapsack / Projection Functions**
Used to **optimize driver allocation** while keeping constraints (e.g., limited idle drivers). Important for **LP-cA2C** to adjust actions using linear programming ideas.

algorithm/alg_utility.py II

- ▶ **Fully-Connected Layer Function (fc)**
Simple wrapper to build neural network layers.
- ▶ **Sampling Utilities (categorical_sample_split)**
Helps sample discrete actions for each agent.
- ▶ **State/Reward Conversion Helpers**
Maps grid states → node IDs that the RL model understands.
- ▶ **Collision Detection**
Counts action conflicts (two agents moving into each other's grid).
- ▶ **Q-table Summation**
Utility for debugging: checks overall Q-value magnitude.

alg_utility.py provides math tools, distribution functions, sampling, projections, and helper utilities that the RL algorithms need to make decisions, update networks, and enforce constraints.

algorithm/IDQN.py I

Purpose: Implements **Independent Deep Q-Learning** for ride-sharing grids — each grid learns its own Q-value without coordination.

Key Ideas (Very Simple):

- ▶ **Estimator (DQN Network)**

3-layer neural network → outputs $Q(s, a)$ for each possible move. Learns using **Bellman update**: $Q \leftarrow r + \gamma \max(Q_{\text{next}})$.

- ▶ **Action Selection**

Uses **valid neighbor mask** (can only move to connected grids). ϵ -greedy: mostly best action, sometimes explore.

Converts Q-values → number of drivers moving to neighbors.

- ▶ **State Processor**

Converts full city state → **per-grid state vectors**. Adds time one-hot, driver counts, order counts. Computes rewards & next states for training.

algorithm/IDQN.py II

- ▶ **Replay Memory**

Stores experience (s, a, r, s') . Samples random mini-batches for stable training.

- ▶ **ModelParameterCopier**

Copies online network → target network for stable Q-learning.

IDQN.py trains a separate DQN for each grid to decide where its drivers should go next, without any coordination between grids.

algorithm/cDQN.py I

Purpose: Implements **Contextual Deep Q-Network (cDQN)** — *one shared value network for all grids*, coordinating driver movement based on context.

Key Ideas (Very Simple):

- ▶ **One Shared Value Network**

Instead of 504 separate DQNs (IDQN), cDQN uses **one model** to value all grid states. This makes learning **faster and more coordinated**.

- ▶ **Context = Number of Idle Drivers**

If a grid has more idle drivers, it can take more actions. Action distribution based on value of neighbor grids + context.

- ▶ **Expected SARSA Update**

Computes targets using **probability-weighted next Q-values**, not max Q.

- ▶ **Neighbor Masking**

Each grid can only move drivers to **its connected neighbors**.

- ▶ **Action Generation**

For each grid: look at Q-values of neighbor grids → pick best neighbor (ϵ -greedy) → split idle drivers among neighbors.

- ▶ **Replay Memory + Target Network**

Stores past experiences and stabilizes training using parameter copying.

cDQN.py learns a single global value function that directs driver movements across all grids, using context and neighbor structure to coordinate actions better than IDQN.

cA2C.py |

Purpose: Smartly dispatch drivers across city grids using **A2C (Actor–Critic)** with a **central coordinator**.

Key Idea:

- ▶ One centralized Value Network (Critic) – learns how “good” each grid state is.
- ▶ One centralized Policy Network (Actor) – learns where to send drivers next.
- ▶ Uses **Advantage = (reward + future value – current value)** to update policy.

How it Works:

1. **State Processor** – Converts city-wide state → per-grid features (drivers, orders, time, grid ID).

cA2C.py II

2. **Value Network (Critic)** – 3-layer MLP → outputs $V(s)$ for each grid. Trained with TD target:
 $target = reward + \gamma V(next_state)$.
3. **Policy Network (Actor)** – 3-layer MLP → outputs action probabilities. Invalid actions removed using neighbor masks. Trains using $policy_loss = -\log \pi(a|s) * advantage$.
4. **Central Advantage Calculation** – Critic computes future value → advantage for each move.
5. **Replay Buffers** – Stores experiences for policy and value updates.

Actor and Critic look at **all grids together**, not independently (unlike IDQN). Ensures **better global coordination** across the city.

`run/run_baseline_nopolicy.py` |

Purpose: Runs the simulator **without any RL policy**. Drivers do **not move**. Used as a **baseline** to compare with IDQN, cDQN, cA2C.

What It Does:

- ▶ Loads real data (orders, drivers, grid map).
- ▶ Creates the CityReal simulator (full-day, 144 steps).
- ▶ Runs 10 episodes with **empty actions** → no rebalancing.
- ▶ Environment handles orders: drivers only serve orders in their own grid.
- ▶ Collects metrics: reward/GMV, response rate, idle drivers, order count.
- ▶ Saves outputs to `results.pkl`.

The baseline runs the city simulation with **zero movement**, giving the **worst-case benchmark** for evaluating all RL algorithms.

run/run_IDQN.py |

Purpose: Train **Independent DQN (IDQN)** where each grid acts as an **independent agent** with its own Q-value decision.

What It Does:

- ▶ Loads real simulator data (orders, drivers, grid map).
- ▶ Initializes IDQN: one neural network, but decisions are **per-grid and uncoordinated**.
- ▶ Runs 25 training episodes:
 - ▶ Reset city for a full day (144 steps).
 - ▶ Convert city state → per-grid state features.
 - ▶ For each step: choose grid actions via Q-network, execute in simulator, compute reward, store experience, train from replay, update target network.
- ▶ Records metrics: GMV, response rate, dispatched drivers, conflicts.
- ▶ Saves logs and model checkpoints.

run/run_IDQN.py II

Trains separate Q-learning agents on each grid to move drivers without any coordination.

run/run_cDQN.py |

Purpose: Train **Contextual DQN (cDQN)** where **all grids share one Q-network** and decisions use **city-level context**.

What It Does:

- ▶ Loads real city data (orders, drivers, grid map).
- ▶ Initializes cDQN: one shared Q-network using a **context vector**.
- ▶ Runs 25 training episodes:
 - ▶ Reset city for a full day.
 - ▶ Build state = grid features + city context.
 - ▶ At each step: network outputs coordinated driver movement, simulator updates state, reward stored in replay memory.
- ▶ After each episode: large training phase (4000 updates), target network update, save model.
- ▶ Tracks metrics: GMV, response rate, conflicts, dispatched drivers.

run/run_cDQN.py II

Trains one shared Q-network that coordinates all grids using city context for better global rebalancing.

run/run_cA2C.py |

Purpose: Train **Contextual A2C (cA2C)** where each grid uses an **Actor** to choose actions and a **Critic** to evaluate them, both conditioned on **city context**.

What It Does:

- ▶ Loads real city data (orders, drivers, grid map).
- ▶ Initializes Actor network (driver movement) and Critic network (state value).
- ▶ Runs 25 training days:
 - ▶ At each step: Actor outputs actions; simulator applies them; reward + next state collected.
 - ▶ Stores data in two buffers: Critic replay (value targets) and Policy replay (advantages).
- ▶ After each episode: **4000 Critic updates + 4000 Actor updates.**

run/run_cA2C.py II

- ▶ Saves model and logs performance (GMV, response rate, conflicts, dispatched drivers).

Trains an Actor–Critic system that uses contextual city information to coordinate driver dispatch across all grids.

simulator/envs.py (CityReal) |

Purpose: Simulates a real city for ride-hailing dispatch using grids, drivers, orders, and time-based distributions.

Core Idea: CityReal is a **time-driven simulation** that models order arrivals, driver movement, ride durations, and prices, enabling RL algorithms (cDQN, cA2C, IDQN) to learn dispatching.

What CityReal Does:

- ▶ Represents city as $M \times N$ grid nodes (valid/invalid areas).
- ▶ Tracks drivers (idle, offline, online).
- ▶ Generates orders using historical distributions (Poisson, Gaussian).
- ▶ Simulates order durations, destinations, and prices.
- ▶ Controls online/offline drivers based on real data.
- ▶ Provides observations for RL (grid-wise idle drivers and active orders).

simulator/envs.py (CityReal) II

- ▶ Steps forward in time (10-minute intervals) for a full day.

Key Components:

- ▶ `construct_node_real` → creates nodes in valid areas.
- ▶ `utility_bootstrap_oneday_order` → samples one day's orders.
- ▶ `initial_order_random` / `step_generate_order_real` → generates timestep orders.
- ▶ `utility_add_driver_real` / `set_drivers_offline` → controls driver counts and locations.
- ▶ `get_observation` → returns 2-channel state for RL (drivers + orders).
- ▶ `reset` / `reset_clean` → initializes city for simulation.

Provides RL models a realistic, data-driven city environment to learn dispatching strategies matching real-world conditions.

simulator/objects.py |

Purpose:

- ▶ Simulates a ride-hailing environment with **Nodes**, **Drivers**, **Orders**, and **Order Distributions**.
- ▶ Handles order generation, driver status update, and simple order–driver assignment.

1. Distribution (Abstract Class)

- ▶ Defines interface for sampling random order counts.

PoissonDistribution

- ▶ Samples number of orders from $\text{Poisson}(\lambda)$.

GaussianDistribution

- ▶ Samples number of orders from a $\text{Normal}(\mu, \sigma)$.

2. Node

- ▶ Represents a grid-cell (location) in the city.

simulator/objects.py II

- ▶ Stores:
 - ▶ Neighboring nodes
 - ▶ Active orders
 - ▶ Drivers (idle/offline)
- ▶ Generates orders (random or real distribution).
- ▶ Assigns orders to drivers within node or neighboring nodes.
- ▶ Updates/removes completed or expired orders.

3. Driver

- ▶ Represents a driver in the system.
- ▶ Tracks online/offline and serving/not-serving status.
- ▶ Takes orders and moves to destination when completed.
- ▶ Updates its internal clock and state each timestep.

4. Order

- ▶ Represents a customer request.

simulator/objects.py III

- ▶ Stores origin node, destination node, start time, duration, price.
- ▶ Tracks assigned time and waiting time.

simulator/utilities.py |

Purpose:

- ▶ Provides helper utilities for the simulation: file handling, time range generation, grid indexing, and neighbor search.

1. File & Time Utilities

- ▶ `datetime_range`: Generates timestamps between start and end with a fixed step.
- ▶ `mkdir_p`: Safe directory creation (like Linux `mkdir -p`).

2. Grid Index Conversion

- ▶ `ids_2dto1d`: Converts (row, col) to 1D index in an $M \times N$ grid.
- ▶ `ids_1dto2d`: Inverse mapping from 1D index back to (row, col).

3. Neighbor Computation

simulator/utilities.py //

- ▶ `get_neighbor_list`: Returns actual neighbor **Node objects** for hexagonal (6-side) or square (4-side) grid.
- ▶ `get_neighbor_index`: Returns neighbor coordinates around a cell.
- ▶ `get_layers_neighbors`: Finds neighbors layer-by-layer (1-hop, 2-hop, ..., L hops).

4. Driver Node State Utilities

- ▶ `get_driver_status`: Builds a matrix of idle drivers over the grid.
- ▶ `debug_print_drivers`: Prints debugging info of all drivers in a node.

tests/run_example.py |

Purpose

- ▶ Runs a small-scale simulation using the CityReal environment.
- ▶ Tests how orders, drivers, and grid locations interact over 144 time steps.
- ▶ Prints order response rate and driver statistics at each step.

Classes

- ▶ **CityReal:** Simulates city-level ride-hailing operations.
Manages orders, drivers, grid states, and time evolution.
- ▶ **Grid/Mapping Utilities (from simulator.envs):** Convert city map into grid representation; support movement, neighborhood, and indexing.
- ▶ **Driver State Manager:** Tracks idle, active, and on/off drivers during simulation.

tests/run_example.py II

- ▶ **Order Generator:** Produces synthetic orders based on duration, price distribution, and grid locations.

Core Simulation Flow

1. Load a 3×3 grid map and generate:
 - ▶ order distribution per grid,
 - ▶ idle driver distribution,
 - ▶ real orders for each time step.
2. Create the CityReal environment with these inputs.
3. Reset the environment and run 144 iterations:
 - ▶ At each step call `env.step()` with empty dispatch action.
 - ▶ Environment internally:
 - 3.1 Generates new orders.
 - 3.2 Matches drivers (if possible).
 - 3.3 Updates city state and metrics.
4. Print order response rate and driver statistics at each step.
5. At the end, compute mean order response rate.

tests/run_example.py III

Outcome

- ▶ Demonstrates basic driver-order dynamics.
- ▶ Shows how response rate changes under no-dispatch policy.

Baseline Simulation: run_example.py I

Purpose: Evaluate city simulator **without any RL policy** (drivers do not move). Serves as **benchmark**.

Simulation Setup:

- ▶ City grid: 3×3 (mapped_matrix_int)
- ▶ Time steps: 144 (full day)
- ▶ Idle driver distribution: 2 per grid per timestep
- ▶ Orders: predefined set per timestep
- ▶ No rebalancing actions (dispatch_action = [])

Metrics Collected:


- ▶ Order Response Rate
- ▶ Idle Drivers
- ▶ Total Drivers

Baseline Simulation: run_example.py II

Results:

- ▶ Order response rate: **avg** ≈ 0.998 over 144 steps
- ▶ Idle drivers fluctuate as drivers serve orders within grids
- ▶ Total drivers remain constant (no new drivers added)

Visualization: Plotted **Order Response Rate**, **Idle Drivers**, and **Total Drivers** over time for quick analysis.



Frame Title




Figure: Screenshot of result from tests/run_example.py